

Load

big data technology


week 3:

map-reduce and programming assignment

week 4:

distributed file-systems, databases, and trends



parallel computing



speedup, $S = T_1 / T_p$, time
with p processors vs with one

efficiency, $E = T_1 / p T_p$

scalable algorithm – E
increasing function of n/p
where n is ‘problem size’



parallel programming paradigms

shared memory

– partition work

$F(w_p)$:

shared a

lock($a[i]$)

work(w_p)

unlock($a[i]$)

message passing

– partition data

$F(p)$:


$a_p = a[p \dots p+(n/p)-1]$

work(w)

exchange data(a_p)

shared + partition data; message-passing + partition work also possible
map-reduce: message-passing, data-parallel, pipelined work, higher level

map-reduce



mappers:

take in k_1, v_1 pairs

emit k_2, v_2 pairs

$k_2, v_2 \leftarrow \text{map}(k_1, v_1)$

reducers:

receive all pairs for some k_2

combine these in some manner

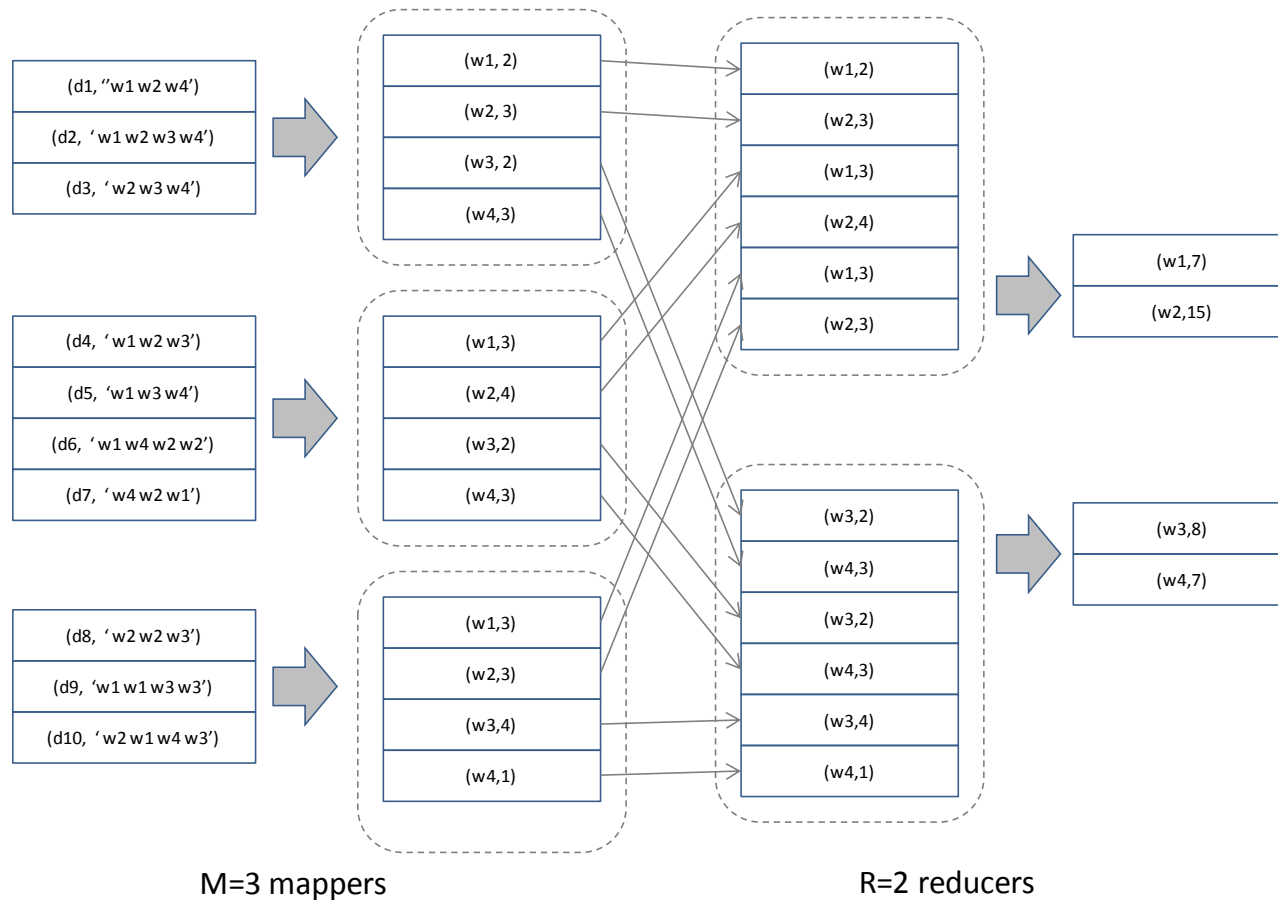
$k_2, f_r(\dots v_2 \dots) \leftarrow \text{reduce}(k_2, [\dots v_2 \dots])$

map-reduce platform responsible for *routing* pairs to reducers
map-reduce reads data and *writes* fresh data; is a *batch* process

map-reduce

Map: $(d_k, 'w_1 \dots w'_n) \rightarrow [(w_i, c_i)]$ document -> word-count pairs

Reduce: $(w_i, [c_i]) \rightarrow (w_i, \sum_i c_i)$ word, count-list -> word-count-total

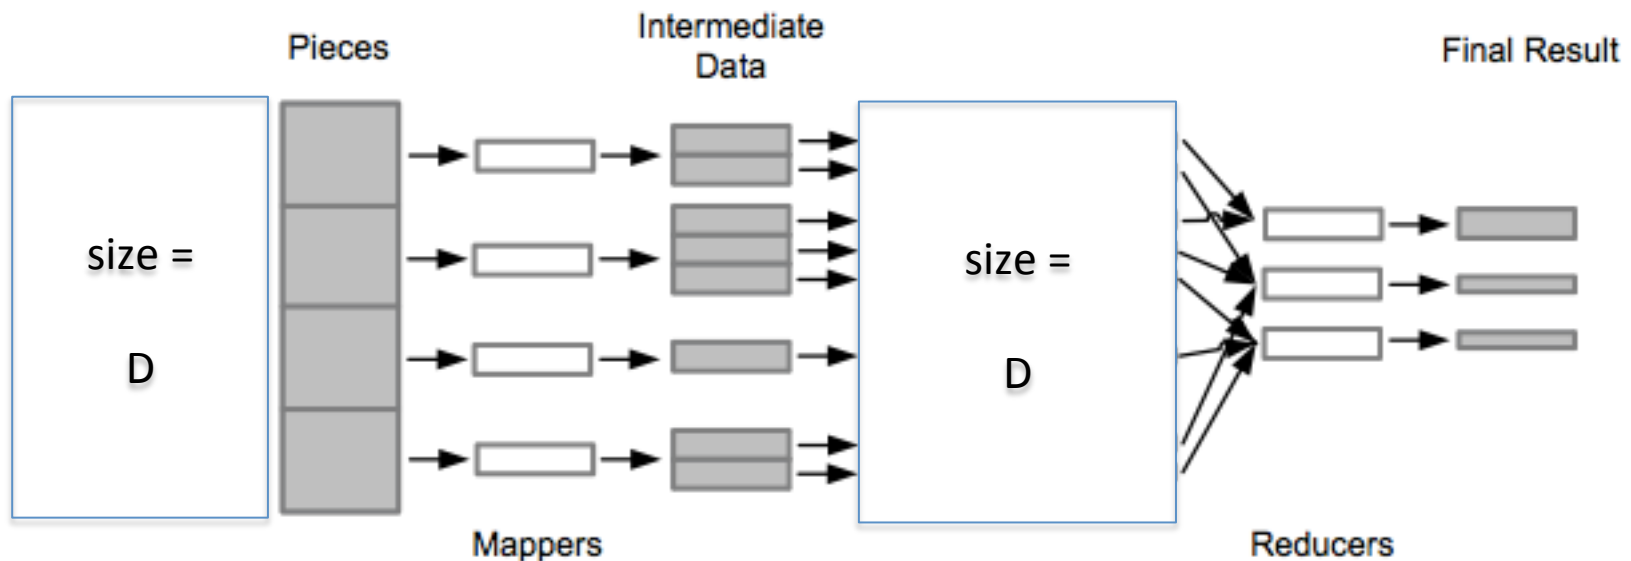


map, reduce ... also 'combine'

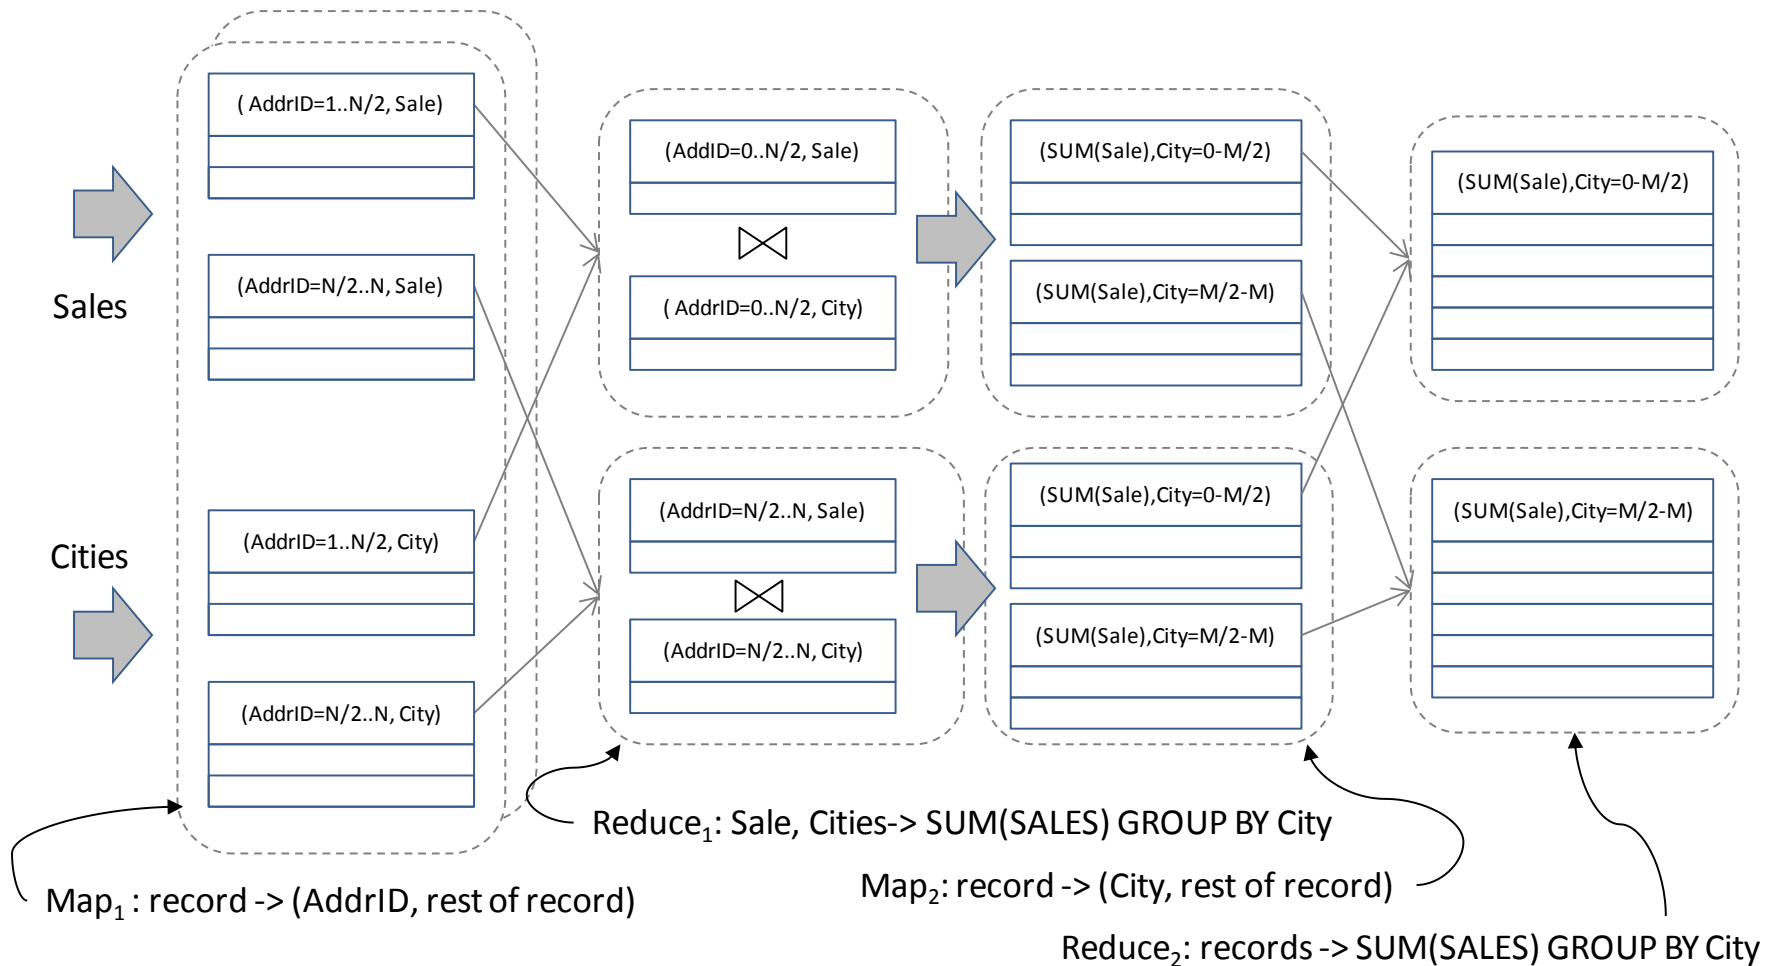
how much data is produced by map?

each word is emitted multiple times!

combiner : sum up word-counts per mapper before emitting



database join using map-reduce



real-world example

lots of data ...

paper, author, contents

million such papers, million authors, millions of possible terms ('phrases' occurring in contents)

problems:

top 10 terms for each author; top 10 authors per term...

'database' person's solution

Q = select id, word, author
from P where in(w,content)

select count(), word, author
from Q group by word

id (paper-id)	content	author



id	word	author




wc	word	author

P million

Q

trillions (million x million)!

top-k words per author in map-reduce




map:
emit word, author

reduce:
reduce-key = word+author
reduce-function = count

suffers from same problem – trillion combinations!

– map-reduce alone is not enough – *approach needs to change!*

top-k words per author in map-reduce



map:
emit author, contents

reduce:
reduce-key = author
reduce-function = **F()**


F(): for each author:

scan all inputs and compute word-counts .. insert into w
sort w, output the top k, delete w and reinitialize to []

look, listen examples in map-reduce

- indexing
- locality-sensitive hashing – how to assemble
- likelihoods – for Bayesian classification
- likelihood ratio – do you need parallelism?
- TF-IDF - HW
- joint probabilities - HW

indexing in map-reduce




map:
produce a partial index
i.e. emit $w \rightarrow$ postings-list

reduce:
reduce-key = word
merge partial indexes
i.e. merge postings per word

what about sorting by either document-id, or page-rank etc. ?

LSH in map-reduce




map:
emit doc-id, k hash-values

reduce:
reduce-key = hashes
emit doc-pairs for each key

will a document-pair be emitted by more than one reducer?

likelihoods in map-reduce



map:

emit counts (f, yes), (f, no)

reduce:


reduce-key = features

sum the counts, divide by N_f

emit the log-likelihoods

once we have the log-likelihoods for each features, do we need parallelism for testing new documents using naïve Bayes?

parallel efficiency of map-reduce



σD data (post map), P processors – mappers + reducers
assume wD is the useful work needs to be done.

Overheads: $\frac{\sigma D}{P}$ intermediate data is written by each mapper

the time for transmitting it to P reducers: $\frac{\sigma D}{P^2} \times P = \frac{\sigma D}{P}$

$$\epsilon_{MR} = \frac{wD}{P(\frac{wD}{P} + 2c\frac{\sigma D}{P})} = \frac{1}{1 + \frac{2c}{w}\sigma}$$

scalable: efficiency approaches 1 as useful work per data-item w grows,
independent of P

parallel-efficiency of MR word-counting

n documents, m words, occurring f times per document on average, so $D = nmf$

the map phase produces mP partial counts,

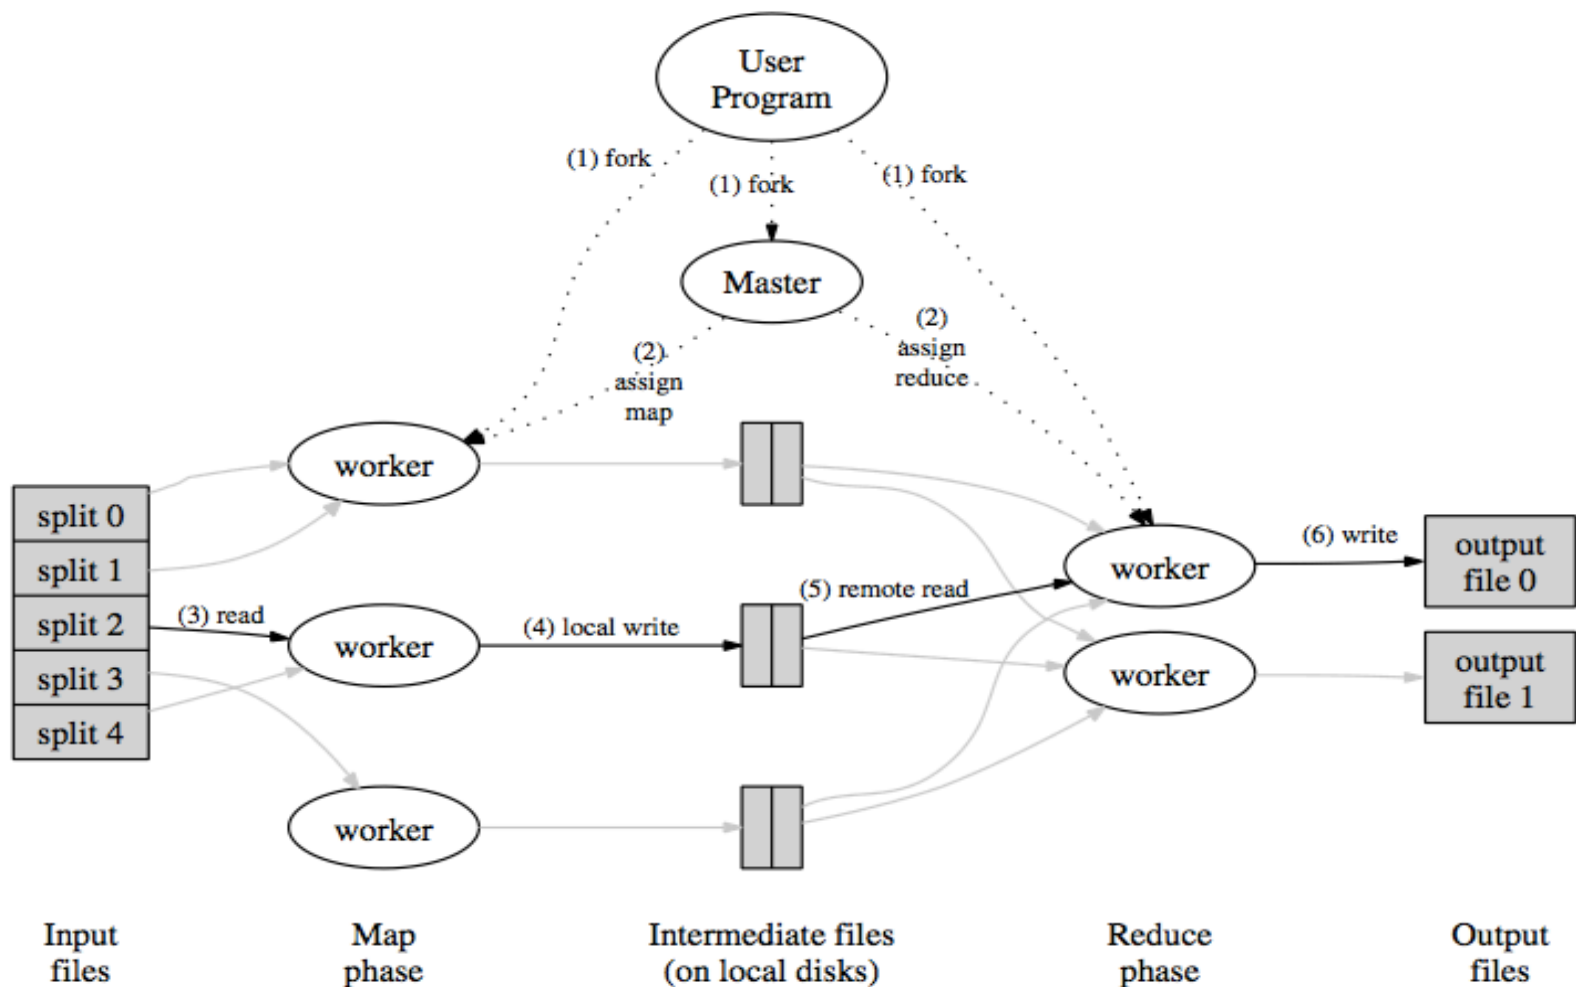
$$\sigma = \frac{mP}{nmf} = \frac{P}{nf}$$

and

$$\varepsilon_{MR} = \frac{1}{1 + \frac{2cP}{wnf}} = \frac{1}{1 + \frac{2P}{nf}}$$

now, scalability is evident as $\frac{n}{p} \rightarrow \infty$

inside map-reduce



recap and preview

parallel computing

map-reduce, applications, internals

Next week:

distributed file systems

distributed (no-SQL) databases

emerging trends